


# Operating System Software to support Virtual Memory

Chapter 8.2, Livro do William Stallings  
SISTEMAS OPERATIVOS, 1º Semestre, 2004-2005

## A VM Larger Than Physical Memory



## Operating System Policies for Virtual Memory

- **Fetch Policy**
  - Demand
  - Prepaging
- **VM: Files and Processes**
- **Placement Policy**
- **Replacement Policy**
  - Basic Algorithms:**
    - Optimal
    - LRU
    - FIFO
    - Clock
    - Page Buffering
- **Resident Set Management**
  - Resident Set Size
    - Fixed
    - Variable
  - Replacement Scope
    - Global
    - Local
- **Cleaning Policy**
  - Demand
  - Precleaning
- **Load Control**
  - Degree of multiprogramming


## Fetch Policy

- **Fetch Policy**
  - Determines when a page should be brought into memory.
  - **Demand paging** only brings pages into main memory when a reference is made to a location on the page.
    - Many page faults when process first started.
  - **Prepaging** brings in more pages than needed
    - More efficient to bring in pages that reside contiguously on the disk.

## Demand Paging

- Bring a page into memory only when it is needed.
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users
- Page is needed  $\Rightarrow$  reference to it
  - invalid reference  $\Rightarrow$  abort
  - not-in-memory  $\Rightarrow$  bring to memory

## Demand Paging



## Performance of Demand Paging

- Page Fault Rate  $0 \leq p \leq 1.0$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault

- **Effective Access Time (EAT)**

$$\text{EAT} = (1 - p) \times \text{memory access} + p \text{ [page fault overhead]}$$

[swap out + swap in + OS overhead]



## Example

- Average page fault overhead: 25 milliseconds
- Memory access time: 100 nanoseconds

- $\text{EAT} = (1 - p) \times (100) + p (25,000,000)$   
 $= 100 + 24,999,900 \times p$

EAT: is directly proportional to the page-fault rate.

If one access out of 1000 causes a page-fault,

$p = 0.001 \rightarrow \text{EAT} = 25 \text{ microseconds}$ .

The computer will be slowed down by a factor of 250!...

## Another Example

- Memory access time = 40 ns ( $4 \times 10^{-8}$ s)
  - 50% of time page being replaced has been modified and so must be written back to disk.
  - Each page fault: swap in page, 50% written back.
  - Swap Page Time (IN or OUT) = 10 msec =  $1 \times 10^{-2}$ s
- EAT       $= (1 - p) \times 4 \times 10^{-8}\text{s} + p (1.5 \times 10^{-2}\text{s} + \text{OS\_Overhead})$   
          $= 4 \times 10^{-8}\text{s} + 0.01499996p\text{s}$   
          $= (40 + 14999960p)\text{ns}$

## Virtual Memory: Processes and Files

## Virtual Memory: Processes and Files

- Virtual memory allows other benefits during process creation and when using files:
  - Copy-on-Write
  - Memory-Mapped Files

## Copy-on-Write

- Copy-on-Write allows both parent and child processes to initially *share* the same pages in memory.

If either process modifies a shared page, only then is the page copied.
- Copy-on-Write allows more efficient process creation as only modified pages are copied.

## Copy-on-Write vs vfork()

- **Copy-on-Write** is used in many operating systems including Linux, Windows 2000/XP, and Solaris.
- Some systems use a version of fork() - **vfork()**.
- **vfork()** is different than copy-on-write.
- **vfork()** suspends the parent until the child has called **exec()** or has exited.
- **vfork()** is intended for use when the child calls **exec()** immediately after being forked. If the child does anything to the address space of the parent before calling **exec()**, the changes will be reflected in the parent as well.

## Memory-Mapped Files

- Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping a disk block to a page in memory**.
- A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- Simplifies file access by treating file I/O through memory rather than **read() write()** system calls.
- Allows several processes to map the same file allowing the pages in memory to be shared.

## Memory-Mapped Files: API

- **Mapping a file**

```
void *mmap(void *addr, size_t len, int prot,  
int flags, int fildes, off_t off);
```

- **Unmapping the file**


```
int munmap(caddr_t addr, size_t len);
```

## Memory-Mapped Files: Example

```
int main(int argc, char *argv[])  
{  
    int fd, offset;  
    char *data;  
    struct stat sbuf;  
  
    if ((fd = open("xpto", O_RDONLY)) == -1) {  
        perror("open");  
        exit(1);  
    }  
  
    if ((data = mmap((caddr_t)0, sbuf.st_size, PROT_READ, MAP_SHARED, fd, 0)) == (caddr_t)(-1)){  
        perror("mmap");  
        exit(1);  
    }  
  
    offset = 10; // vai buscar o 10º byte do ficheiro...  
    printf("byte at offset %d is '%c'\n", offset, data[offset]);  
    return 0;  
}
```



## Memory-Mapped Files




## Replacement Policy

# Replacement Policy

- **Placement Policy**

- Which page is replaced?
- Use *dirty (modify) bit* to reduce overhead of page transfers – only modified pages written back to disk.
- Page removed should be the page least likely to be referenced in the near future.
- Most policies predict the future behavior on the basis of past behavior.


## Need for Page Replacement



# Basic Page Replacement

1. Find the location of the desired page on disk.
2. Find a free frame:
  - If there is a free frame, use it.
  - If there is no free frame, use a page replacement algorithm to select a *victim frame*.
  - If victim is *dirty* write back to disk
3. Read the desired page into the (newly) free frame.  
Update the page and frame tables.
4. Restart the process.

## Example




## Page Locking and the Replacement Policy

- **Frame Locking**

- If frame is locked, it may not be replaced
- Kernel of the operating system
- Control structures
- I/O buffers (Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm).
- Associate a **lock-bit** with each frame

## Frames Used For I/O Must Be In Memory




## **Page-Replacement Algorithms**

### **Page-Replacement Algorithms**

- Goal: get the lowest page-fault rate
- Evaluate algorithm by running on given string of memory references and compute the number of page faults
- Example of a reference string:  
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5


## Graph of Page Faults Versus The Number of Frames



## Replacement Algorithm: FIFO

- **First-in, first-out (FIFO)**
  - Treats page frames allocated to a process as a circular buffer.
  - Pages are removed in round-robin style.
  - Simplest replacement policy to implement.
  - Page that has been in memory the longest is replaced.
  - These pages may be needed again very soon.

# FIFO Page Replacement



# FIFO Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (for *all* processes: 3 pages in memory at a time)

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults


- 4 frames

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults

- FIFO replacement bad, can lead to – **Belady's Anomaly**
  - more frames  $\Rightarrow$  *more* page faults !!

## Belady's Anomaly




## Replacement Algorithm: LRU

- **Least Recently Used (LRU)**
  - Replaces the page that has not been referenced for the longest time.
  - By the principle of locality, this should be the page least likely to be referenced in the near future.
  - Each page could be tagged with the time of last reference. This would require a great deal of overhead.



# LRU Page Replacement



# LRU Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5


1	5	
2		8 page-faults
3	5	4
4	3	

- Counter implementation
  - every page table entry has timestamp; every time page referenced through entry, copy clock (counter) into timestamp
  - when page needs to be replaced, look at timestamp to choose

## Optimal Replacement Algorithm

- **Optimal policy**
  - Selects for replacement that page for which the time to the next reference is the longest.
  - Impossible to have perfect knowledge of future events.

## Optimal Page Replacement



## Optimal Algorithm



- Replace page not to be used for longest time in future
- 4 frames example: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	4	6 page faults
2		
3		
4	5	

- How do you know which frame to replace?
- Useful as standard for comparing realistic algorithms

## Clock Replacement Algorithm


- **Clock Policy**
  - Additional bit called a **use bit**.
  - When a page is first loaded in memory, the **use bit** is set to 0.
  - When the page is referenced, the **use bit** is set to 1.
  - When it is time to replace a page, the first frame encountered with the **use bit** set to 0 is replaced.
  - During the search for replacement, each **use bit** set to 1 is changed to 0.



## Page Replacement Algorithms

Page address stream	2	3	2	1	5	2	4	5	3	2	5	2																																				
OPT	<table><tr><td>2</td></tr><tr><td></td></tr><tr><td></td></tr></table> F	2			<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table> F	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table> F	2	3	1	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table> F	2	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	2	3	5	<table><tr><td>4</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table> F	4	3	5	<table><tr><td>4</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	4	3	5	<table><tr><td>4</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	4	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table> F	2	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	2	3	5	<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>5</td></tr></table>	2	3	5
2																																																
2																																																
3																																																
2																																																
3																																																
2																																																
3																																																
1																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
LRU	<table><tr><td>2</td></tr><tr><td></td></tr><tr><td></td></tr></table> F	2			<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table> F	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table> F	2	3	1	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>1</td></tr></table> F	2	5	1	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>1</td></tr></table>	2	5	1	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table> F	2	5	4	<table><tr><td>2</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table>	2	5	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table> F	3	5	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table> F	3	5	2	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table>	3	5	2	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table>	3	5	2
2																																																
2																																																
3																																																
2																																																
3																																																
2																																																
3																																																
1																																																
2																																																
5																																																
1																																																
2																																																
5																																																
1																																																
2																																																
5																																																
4																																																
2																																																
5																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
3																																																
5																																																
2																																																
FIFO	<table><tr><td>2</td></tr><tr><td></td></tr><tr><td></td></tr></table> F	2			<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table> F	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td></td></tr></table>	2	3		<table><tr><td>2</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table> F	2	3	1	<table><tr><td>5</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table> F	5	3	1	<table><tr><td>5</td></tr><tr><td>2</td></tr><tr><td>1</td></tr></table> F	5	2	1	<table><tr><td>5</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table> F	5	2	4	<table><tr><td>5</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table>	5	2	4	<table><tr><td>3</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table> F	3	2	4	<table><tr><td>3</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table>	3	2	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>4</td></tr></table> F	3	5	4	<table><tr><td>3</td></tr><tr><td>5</td></tr><tr><td>2</td></tr></table> F	3	5	2
2																																																
2																																																
3																																																
2																																																
3																																																
2																																																
3																																																
1																																																
5																																																
3																																																
1																																																
5																																																
2																																																
1																																																
5																																																
2																																																
4																																																
5																																																
2																																																
4																																																
3																																																
2																																																
4																																																
3																																																
2																																																
4																																																
3																																																
5																																																
4																																																
3																																																
5																																																
2																																																
CLOCK	<table><tr><td>2*</td></tr><tr><td></td></tr><tr><td></td></tr></table> E	2*			<table><tr><td>2*</td></tr><tr><td>3*</td></tr><tr><td></td></tr></table> F	2*	3*		<table><tr><td>2*</td></tr><tr><td>3*</td></tr><tr><td></td></tr></table> →	2*	3*		<table><tr><td>2*</td></tr><tr><td>3*</td></tr><tr><td>1*</td></tr></table> F	2*	3*	1*	<table><tr><td>5*</td></tr><tr><td>3</td></tr><tr><td>1</td></tr></table> F	5*	3	1	<table><tr><td>5*</td></tr><tr><td>2*</td></tr><tr><td>1</td></tr></table> F	5*	2*	1	<table><tr><td>5*</td></tr><tr><td>2*</td></tr><tr><td>4*</td></tr></table> F	5*	2*	4*	<table><tr><td>5*</td></tr><tr><td>2*</td></tr><tr><td>4*</td></tr></table> →	5*	2*	4*	<table><tr><td>3*</td></tr><tr><td>2</td></tr><tr><td>4</td></tr></table> F	3*	2	4	<table><tr><td>3*</td></tr><tr><td>2*</td></tr><tr><td>4</td></tr></table> →	3*	2*	4	<table><tr><td>3*</td></tr><tr><td>2</td></tr><tr><td>5*</td></tr></table> F	3*	2	5*	<table><tr><td>3*</td></tr><tr><td>2*</td></tr><tr><td>5*</td></tr></table> →	3*	2*	5*
2*																																																
2*																																																
3*																																																
2*																																																
3*																																																
2*																																																
3*																																																
1*																																																
5*																																																
3																																																
1																																																
5*																																																
2*																																																
1																																																
5*																																																
2*																																																
4*																																																
5*																																																
2*																																																
4*																																																
3*																																																
2																																																
4																																																
3*																																																
2*																																																
4																																																
3*																																																
2																																																
5*																																																
3*																																																
2*																																																
5*																																																

## Local Page Replacement Algorithms



**goto QUIZ#5...**

## **Resident Set Size**

- **Fixed-allocation**
  - gives a process a fixed number of pages within which to execute.
  - when a page fault occurs, one of the pages of that process must be replaced.
- **Variable-allocation**
  - number of pages allocated to a process varies over the lifetime of the process.

## Replacement Scope

- **Local replacement policy:**
  - Chooses only among the resident pages of the process that generated the page fault in selecting a page to replace.
- **Global replacement policy:**
  - Considers all unlocked pages in main-memory as candidates for replacement.

## Fixed Location, Local Scope

- A process is running in memory with a fixed number of frames.
- When a page fault occurs the OS must choose which page to replace, among the resident pages of this process.
- With a fixed-allocation policy it is necessary to decide ahead of time the amount of allocation to give to a process.
- Trade-off decision...

## **Variable Allocation, Global Scope**

- Easiest to implement; adopted by many operating systems.
- Operating system keeps list of free frames.
- Free frame is added to resident set of process when a page fault occurs.
- If no free frame, replaces one from another process.

## **Variable Allocation, Local Scope**

- When new process added, allocate number of page frames based on application type, program request, or other criteria.
- When page fault occurs, select page from among the resident set of the process that suffers the fault.
- Reevaluate allocation from time to time.



## Cleaning Policy

- **Demand cleaning**
  - a page is written out only when it has been selected for replacement.
- **Pre-cleaning**
  - pages are written out in batches.


## Cleaning Policy

- Best approach uses **page buffering**
  - Replaced pages are placed in two lists
    - Modified and unmodified.
  - Pages in the modified list are periodically written out in batches.
  - Pages in the unmodified list are either reclaimed if referenced again or lost when its frame is assigned to another page.

## Load Control

- Determines the number of processes that will be resident in main memory.
- Too few processes, many occasions when all processes will be blocked and much time will be spent in swapping.
- Too many processes will lead to: thrashing.

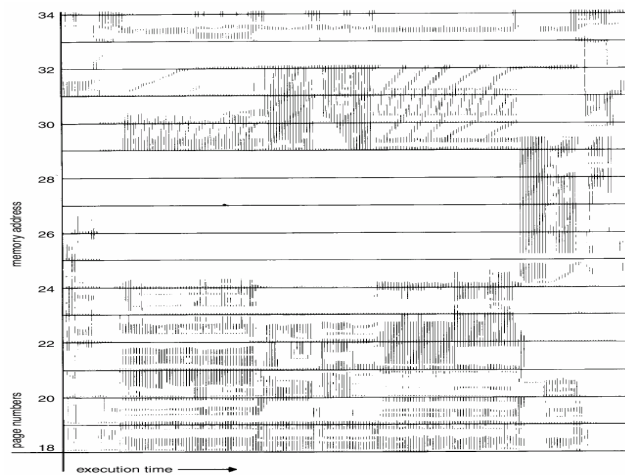
## Multiprogramming Effects



## Why Trashing occurs?

- Why does paging work?  
Locality model
  - Process migrates from one locality to another.
  - Localities may overlap.
- Why does thrashing occur?  
 $\Sigma$  size of locality > total memory size


## Locality in a Memory-Reference Pattern



## Working-Set Model

- $\Delta \equiv$  working-set window  $\equiv$  a fixed number of page references  
Example: 10,000 instruction
- $WSS_i$  (working set of Process  $P_i$ ) = total number of pages referenced in the most recent  $\Delta$  (varies in time)
  - if  $\Delta$  too small will not encompass entire locality.
  - if  $\Delta$  too large will encompass several localities.
  - if  $\Delta = \infty \Rightarrow$  will encompass entire program.
- $D = \sum WSS_i \equiv$  total demand frames
- ( $m$  = total number of page frames)
- if  $D > m \Rightarrow$  **Thrashing**
- Policy: if  $D > m$ , then suspend one of the processes.

## Working Set Model



## Exemplo: dois programas...

```
int A[][] = new int[128][128];  
int i,j;
```

```
// page size = 512 bytes  
// each row is stored in one  
// page
```

**// Programa A**

```
for (j = 0; j < 128; j++)  
  for (i = 0; i < 128; i++)  
    A[i][j] = 0;
```



```
int A[][] = new int[128][128];  
int i,j;
```

```
// page size = 512 bytes  
// each row is stored in one  
// page
```

**// Programa B**

```
for (i = 0; i < 128; i++)  
  for (j = 0; j < 128; j++)  
    A[i][j] = 0;
```

**Quem escreveu o programa A?  
E o programa B?**



## Exemplo: dois programas...

```
int A[][] = new int[128][128];
```

// Programa A

```
for (j = 0; j < 128; j++)  
  for (i = 0; i < 128; i++)  
    A[i][j] = 0;
```



```
int A[][] = new int[128][128];
```

// Programa B

```
for (i = 0; i < 128; i++)  
  for (j = 0; j < 128; j++)  
    A[i][j] = 0;
```

128 x 128 = 16,384 page faults

128 page faults



```
int A[][] = new int[128][128];
```

// Programa A

```
for (j = 0; j < 128; j++)  
  for (i = 0; i < 128; i++)  
    A[i][j] = 0;
```

```
int A[][] = new int[128][128];
```

// Programa B

```
for (i = 0; i < 128; i++)  
  for (j = 0; j < 128; j++)  
    A[i][j] = 0;
```

128 x 128 = 16,384 page faults

128 page faults

## Process Suspension

- **Lowest priority process.**
- **Faulting process**
  - this process does not have its working set in main memory so it will be blocked anyway.
- **Last process activated**
  - this process is least likely to have its working set resident.

## Process Suspension

- **Process with smallest resident set**
  - this process requires the least future effort to reload
- **Largest process**
  - obtains the most free frames
- **Process with the largest remaining execution window.**

# Overview of Virtual Memory Systems:

- Unix and Solaris
- Linux
- Windows 2000

## UNIX and Solaris Memory Management

- Two separate memory management schemes in Unix SVR4 and Solaris
  - Paging System – Allocate page frames
  - Kernel Memory Allocator – Allocate memory for the kernel
- Paging System Data Structures
  - Page Table – One entry for each page of virtual memory for that process
    - Page Frame Number – Physical frame #
    - Age – How long in memory w/o reference
    - Copy on Write – Are two processes sharing this page: after fork(), waiting for exec()
    - Modify – Page modified?
    - Reference – Set when page accessed
    - Valid – Page is in main memory
    - Protect – Are we allowed to write to this page?



## Paging Data Structures

- Disk Block Descriptor
  - Swap Device Number – Logical device
  - Device Block Number – Block location
  - Type of storage – Swap or executable, also indicate if we should clear first
- Page Frame Data Table
  - Page State – Available, in use (on swap device, in executable, in transfer)
  - Reference Count – # processes using page
  - Logical Device – Device holding copy
  - Block Number – Location on device
  - Pfdata pointer – For linked list of pages
- Swap-use Table
  - Reference Count – # entries pointing to a page on a storage device
  - Page/storage unit number – Page ID

## Data Structures

Page frame number	Age	Copy on write	Modify	Reference	Valid	Protect
-------------------	-----	---------------	--------	-----------	-------	---------

(a) Page table entry

Swap device number	Device block number	Type of storage
--------------------	---------------------	-----------------

(b) Disk block descriptor

**Figure 8.22 UNIX SVR4 Memory Management Formats**

## Data Structures

Page state	Reference count	Logical device	Block number	Pfdata pointer
------------	-----------------	----------------	--------------	----------------

(c) Page frame data table entry

Reference count	Page/storage unit number
-----------------	--------------------------

(d) Swap-use table entry


Figure 8.22 UNIX SVR4 Memory Management Formats

## UNIX and Solaris Memory Management

- Page Replacement
  - refinement of the clock policy
- Kernel Memory Allocator
  - most blocks are smaller than a typical page size

## Unix SVR4 Page Replacement

- Clock algorithm variant
- *Fronthand* – Clear Use bits
- *Backhand* – Check Use bits, if use=0 prepare to swap page out
- *Scanrate* – How fast the hands move
  - Faster rate frees pages faster
- *Handspread* – Gap between hands
  - Smaller gap frees pages faster
- System adjusts values based on free memory

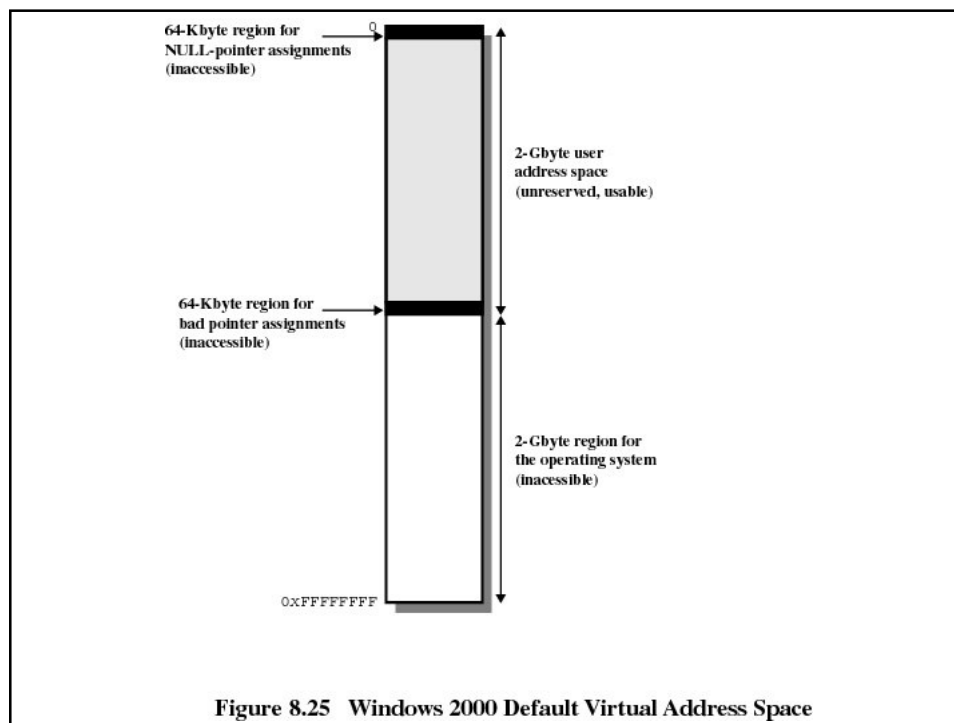


## Linux Memory Management

- Virtual Memory Addressing
  - Supports 3-level page tables
    - Page Directory
      - One page in size (must be in memory)
    - Page Middle Directory
      - Can span multiple pages
      - Will have size=1 on Pentium
    - Page Table
      - Points to individual pages
- Page Allocation
  - Uses a buddy system with 1-32 page block sizes
- Page Replacement
  - Based on clock algorithm
  - Uses age variable
    - Incremented when page is accessed
    - Decrement as it scans memory
    - When age=0, page may be replaced
  - Has effect of least frequently used method
- Kernel Memory Allocation
  - Uses scheme called *slab allocation*
  - Blocks of size 32 through 4080 bytes

# Windows 2000 Memory Management

- Virtual Address Map
  - 00000000 to 00000FFF – Reserved
    - Help catch NULL pointer uses
  - 00001000 to 7FFFFFFF – User space
  - 7FFFFFFF to 7FFFFFFF – Reserved
    - Help catch wild pointers
  - 80000000 to FFFFFFFF – System
- Page States
  - Available – Not currently used
  - Reserved – Set aside, but not counted against memory quota (not in use)
    - No disk swap space allocated yet
    - Process can declare memory that can be quickly allocated when it is needed
  - Committed: space set aside in paging file (in use by the process)



## **Windows 2000 Memory Management**

- Uses variable allocation, local scope.
- When a page fault occurs, a page is selected from the local set of pages.
- If main memory is plentiful, allow the resident set to grow as pages are brought into memory.
- If main memory is scarce, remove less recently accessed pages from the resident set.